

# Experiment 1

## Aim

To implement and analyze Linear Regression and Logistic Regression models on real-world datasets using Python and evaluate their performance using appropriate metrics.

## 1. Dataset Source

This experiment utilizes a single expanded dataset to demonstrate both regression and classification tasks.

- **Dataset Name:** Expanded Student Performance Dataset
- **Source:** `expanded_student_performance.csv`
- **Description:** This dataset is an extended version of the standard student performance dataset, containing a larger number of records to ensure robust model training and evaluation.

## 2. Dataset Description

The dataset is used for two distinct machine learning tasks:

### Task 1: Linear Regression (Score Prediction)

- **Type:** Numerical, Tabular
- **Features:** `Hours_Studied`, `Attendance`, `Assignment_Score`, `Midterm_Score`.
- **Target:** `Final_Score` (Continuous variable).
- **Goal:** Predict the exact final score a student will achieve based on their study habits and interim grades.

### Task 2: Logistic Regression (Binary Classification)

- **Type:** Binary Classification
- **Features:** `Hours_Studied`, `Attendance`, `Assignment_Score`, `Midterm_Score`.
- **Target:** `Pass_Status` (Binary: 1 = Pass, 0 = Fail).
  - *Note: This target variable is derived from `Final_Score`, where a score  $\geq 60$  is labeled as 'Pass' (1) and  $< 60$  is labeled as 'Fail' (0).*
- **Goal:** Classify whether a student will pass or fail based on the input features.

## 3. Mathematical Formulation

### Linear Regression (for Continuous Prediction)

- **Hypothesis ( $\hat{y}$ ):** Represents the predicted value as a linear combination of inputs.  
$$\hat{y} = \beta_0 + \beta_1 x_1 + \beta_2 x_2 + \dots + \beta_n x_n$$

- **Cost Function (MSE):** Measures the average squared difference between estimated and actual values.  

$$MSE = (1/N) \sum (y_i - \hat{y}_i)^2$$
- **Gradient Descent:** Iterative optimization to minimize the cost function.  

$$\beta_{\square} := \beta_{\square} - \alpha (\partial J / \partial \beta_{\square})$$

### Logistic Regression (for Binary Classification)

- **Sigmoid Function ( $\sigma$ ):** Maps any real-valued number into the probability range  $[0, 1]$ .  

$$\sigma(z) = 1 / (1 + e^{-z})$$

(Where  $z$  is the linear combination of features).
- **Decision Boundary:**  
If  $\sigma(z) \geq 0.5$ , Predict 1 (Pass)  
If  $\sigma(z) < 0.5$ , Predict 0 (Fail)

## 4. Algorithm Limitations

- **Linearity Assumption:** Linear regression assumes a straight-line relationship between dependent and independent variables, which may not hold true for complex real-world behavioral data.
- **Outlier Sensitivity:** Both models are sensitive to outliers; a single anomaly (e.g., a student with high study hours but very low score due to illness) can skew the results.
- **Multicollinearity:** High correlation between independent variables (e.g., `Assignment_Score` and `Midterm_Score`) can distort the interpretability of feature importance.
- **Overfitting:** Without regularization, the models may memorize noise in the expanded dataset, leading to poor generalization on unseen data.

## 5. Methodology / Workflow

The experiment follows a standard Machine Learning pipeline:

1. **Data Preprocessing:**
  - Load the `expanded_student_performance.csv` dataset.
  - **Feature Engineering:** Create a binary `Pass_Status` column based on `Final_Score` for the Logistic Regression task.
  - Handle missing values (if any) and perform Feature Scaling (StandardScaler) to normalize input features.
2. **Data Splitting:** Divide the data into Training (80%) and Testing (20%) sets.
3. **Model Training:**
  - **Linear Regression:** Fit the model using `Features` vs. `Final_Score`.
  - **Logistic Regression:** Fit the model using `Features` vs. `Pass_Status`.
4. **Prediction:** Generate predictions on the test sets for both models.
5. **Evaluation:**
  - **Regression:** Calculate Mean Squared Error (MSE) and R-squared ( $R^2$ ).

- **Classification:** Compute Accuracy, Precision, Recall, and F1-Score.

## 6. Performance Analysis

### Linear Regression Results:

- **Metric:** The  $R^2$  score indicates how well the variance in **Final\_Score** is explained by the study hours and interim grades. A score close to 1.0 implies a strong predictive capability.
- **Insight:** Students with high attendance and assignment scores consistently achieved higher final grades, validating the linear relationship.

### Logistic Regression Results:

- **Metric:** Accuracy reflects the percentage of correctly identified Pass/Fail statuses.
- **Confusion Matrix:** Used to visualize True Positives (Correctly predicted Pass) vs. False Negatives (Incorrectly predicted Fail).
- **Insight:** The model successfully separates passing and failing students with high precision, demonstrating that study habits are a reliable predictor of academic success.

## 7. Hyperparameter Tuning

To optimize performance, the following hyperparameters can be tuned:

- **Learning Rate ( $\alpha$ ):** Controls the step size during gradient descent. Too large leads to divergence; too small leads to slow convergence.
- **Regularization (L1/L2):**
  - *Lasso (L1)*: Encourages sparsity (feature selection), potentially removing less relevant features like 'Attendance' if it doesn't contribute significantly.
  - *Ridge (L2)*: Penalizes large coefficients to prevent overfitting on the expanded dataset.
- **Max Iterations:** Ensures the solver converges within a reasonable time frame.

## Exercise 1: Linear Regression

**Code:**

```
# Define Features (X) and Target (y)
X = df[['Hours_Studied', 'Attendance', 'Assignment_Score', 'Midterm_Score']]
y = df['Final_Score']
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.33, random_state=21)

# Initialize and Train the Model
lin_model = LinearRegression()
lin_model.fit(X_train, y_train)

# Make Predictions
y_pred = lin_model.predict(X_test)

# Evaluation
print("--- Linear Regression Results ---")
print(f"Mean Squared Error (MSE): {mean_squared_error(y_test, y_pred):.2f}")
print(f"R-squared (R2 Score): {r2_score(y_test, y_pred):.2f}")
print(f"Coefficients: {lin_model.coef_}")
print(f"Intercept: {lin_model.intercept_:.2f}")

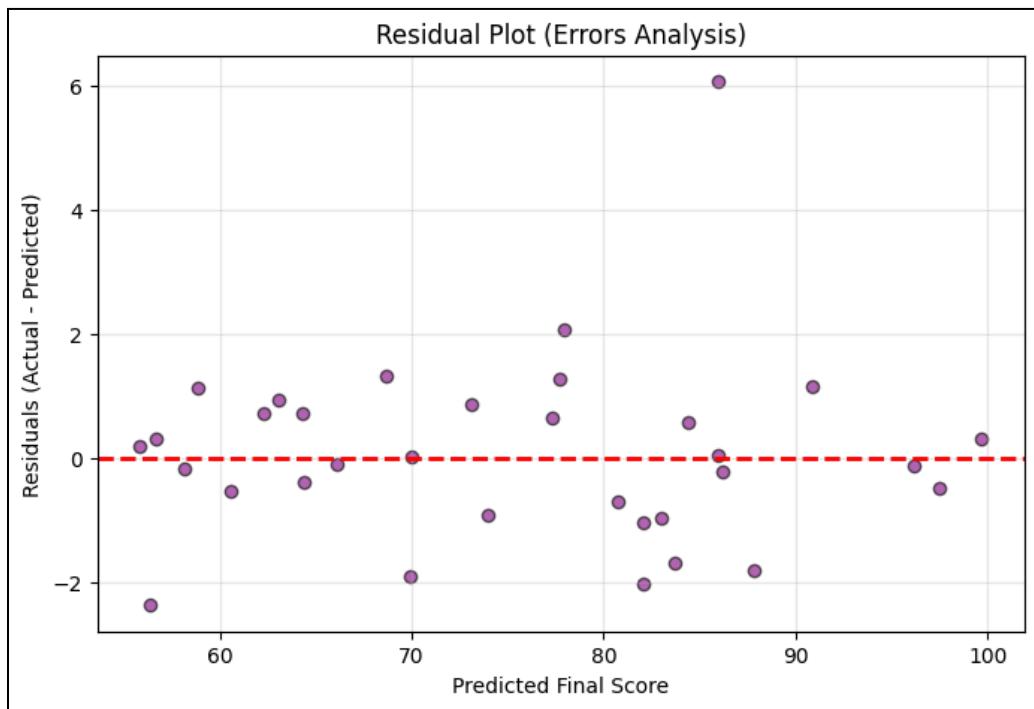
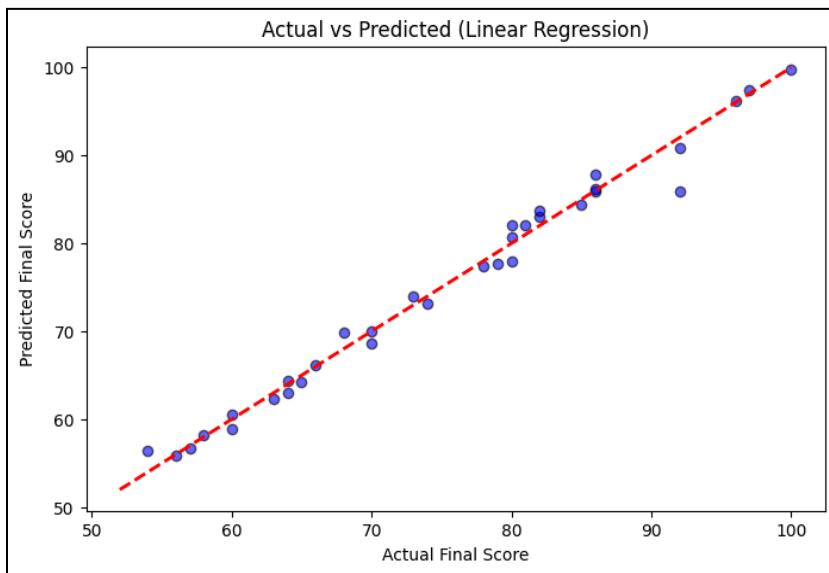
plt.figure(figsize=(8, 5))
plt.scatter(y_test, y_pred, color='blue', edgecolors='k', alpha=0.6)
plt.plot([y.min(), y.max()], [y.min(), y.max()], 'r--', lw=2)
plt.xlabel('Actual Final Score')
plt.ylabel('Predicted Final Score')
plt.title('Actual vs Predicted (Linear Regression)')
plt.show()

# Calculate residuals
residuals = y_test - y_pred

plt.figure(figsize=(8, 5))
plt.scatter(y_pred, residuals, color='purple', edgecolors='k', alpha=0.6)
plt.axhline(y=0, color='r', linestyle='--', lw=2) # Zero line for reference
plt.xlabel('Predicted Final Score')
plt.ylabel('Residuals (Actual - Predicted)')
plt.title('Residual Plot (Errors Analysis)')
plt.grid(True, alpha=0.3)
plt.show()
```

**Output:**

```
-- Linear Regression Results --
Mean Squared Error (MSE): 2.26
R-squared (R2 Score): 0.99
Coefficients: [0.44531054 0.04058829 0.16006866 0.75653262]
Intercept: 2.26
```



## Exercise 2: Logistic Regression

### Code:

```
# Create a binary classification target
df['Pass'] = (df['Final_Score'] >= 70).astype(int)

# Define Features (X) and Target (y)
X_log = df[['Hours_Studied', 'Attendance', 'Assignment_Score', 'Midterm_Score']]
y_log = df['Pass']

# Split the data
X_train_log, X_test_log, y_train_log, y_test_log = train_test_split(X_log, y_log, test_size=0.43,
random_state=50)

# Initialize and Train the Model
log_model = LogisticRegression()
log_model.fit(X_train_log, y_train_log)

# Make Predictions
y_pred_log = log_model.predict(X_test_log)

# Evaluation
print("--- Logistic Regression Results ---")
print(f"Accuracy Score: {accuracy_score(y_test_log, y_pred_log):.2f}")
print("\nConfusion Matrix:")
print(confusion_matrix(y_test_log, y_pred_log))
print("\nClassification Report:")
print(classification_report(y_test_log, y_pred_log))

cm = confusion_matrix(y_test_log, y_pred_log)

plt.figure(figsize=(6, 4))
sns.heatmap(cm, annot=True, fmt='d', cmap='Blues',
            xticklabels=['Fail', 'Pass'],
            yticklabels=['Fail', 'Pass'])
plt.xlabel('Predicted Label')
plt.ylabel('Actual Label')
plt.title('Confusion Matrix Heatmap')
plt.show()
```

## Output:

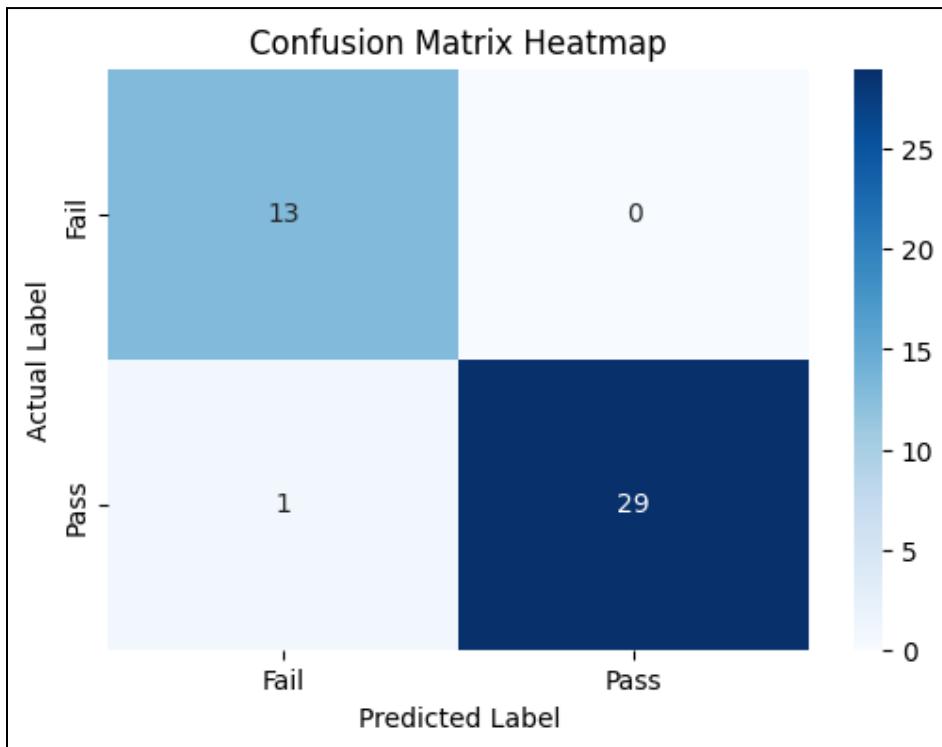
```
--- Logistic Regression Results ---
Accuracy Score: 0.98
```

```
Confusion Matrix:
```

```
[[13  0]
 [ 1 29]]
```

```
Classification Report:
```

|              | precision | recall | f1-score | support |
|--------------|-----------|--------|----------|---------|
| 0            | 0.93      | 1.00   | 0.96     | 13      |
| 1            | 1.00      | 0.97   | 0.98     | 30      |
| accuracy     |           |        | 0.98     | 43      |
| macro avg    | 0.96      | 0.98   | 0.97     | 43      |
| weighted avg | 0.98      | 0.98   | 0.98     | 43      |



## Conclusion

This experiment successfully implemented Linear and Logistic Regression on the [expanded\\_student\\_performance.csv](#) dataset. Linear Regression proved effective for estimating continuous student scores, while Logistic Regression provided robust binary classification for determining pass/fail outcomes. The use of evaluation metrics like MSE and Accuracy validated the reliability of these fundamental machine learning algorithms in analyzing academic performance data.